

## What are Transaction Control Commands?

When you group SQL statements into a transaction, you need a way to **start**, **save**, **cancel**, or **complete** that transaction. These commands are like the buttons you press to control what happens to your group of SQL actions.

Let's break them down one by one in plain language:

### **BEGIN or START TRANSACTION**

#### **What it does:**

This command **marks the beginning** of a transaction. It tells the database, "From this point forward, I'm grouping actions together."

#### **How to explain to students:**

"Imagine you're starting a recipe. You say, 'I'm beginning now, so if I mess up somewhere, I want to be able to throw everything away and start fresh.' That's what **BEGIN** does in SQL. It signals the start of something important that you may want to undo if needed."

### **COMMIT**

#### **What it does:**

This command **makes all the changes permanent**. After a **COMMIT**, the database saves the results of everything done since **BEGIN**.

#### **How to explain to students:**

"Think of **COMMIT** as clicking **Save** on a document. After you commit, your work is stored — it's locked in. Even if the power goes out, the changes stay. You're saying, 'I'm done and happy with everything I've done.'"

### **ROLLBACK**

#### **What it does:**

It **cancels everything** done in the transaction so far — nothing is saved. It returns the database to the way it was when you used **BEGIN**.

#### **How to explain to students:**

"This is like pressing **Undo All** or clicking 'Don't Save' and closing the file. If you made a mistake halfway, **ROLLBACK** lets you erase all the changes and start over."

### **SAVEPOINT**

#### **What it does:**

This creates a **checkpoint** inside a transaction — a point you can return to if something goes wrong after it.

### How to explain to students:

“Imagine you’re writing an essay, and after finishing the introduction, you make a backup copy. That’s a **SAVEPOINT**. If you mess up the next section, you don’t have to throw everything away — you can go back to the backup.”

You give it a name like `SAVEPOINT intro_done;` so you can refer to it later.

## ◆ **ROLLBACK TO SAVEPOINT**

### What it does:

It **undoes part of the transaction** — only the steps that happened after the savepoint.

### How to explain to students:

“This is like saying, ‘I liked what I had before this paragraph — let me go back to that version and fix it from there.’ You don’t erase everything, just the stuff you did after your last savepoint.”

## ◆ **RELEASE SAVEPOINT**

### What it does:

It **removes a savepoint** you no longer need. You can’t roll back to it after it’s released.

### How to explain to students:

“Think of this like deleting that backup copy because you’ve moved on. It helps clean up memory and avoid confusion later on.”



## **Key Points to Emphasize to Students**

- A transaction **starts with** `BEGIN`.
- Once you're sure everything is correct, you use `COMMIT`.
- If there's a mistake, use `ROLLBACK` to cancel.
- You can create **mini-backups** using `SAVEPOINT`.
- You can go back to a **safe point** using `ROLLBACK TO SAVEPOINT`.
- Once done, **clean up** with `RELEASE SAVEPOINT`.

## **4. Autocommit Behavior – Explained for Beginners**



### **What is Autocommit?**

“When you write and run a single SQL statement — like an `INSERT`, `UPDATE`, or `DELETE` — many databases **automatically save the changes** right after that statement finishes. This is called **autocommit**.

It means:

- ➡ You don’t have to say `COMMIT` — the database does it for you behind the scenes.
- ➡ Every statement is treated like a tiny, automatic transaction.”

## Why This Matters

“This seems convenient, but it can be **dangerous** if you’re running multiple steps that need to go together.

Let’s say you want to do these in one group:

1. Subtract money from Account A
2. Add money to Account B

If autocommit is ON, and the first statement runs but the second fails, only half of your task happens. That’s bad — your data is now **incomplete** and **inconsistent**.

You want **both** steps to succeed together or **none at all** — that’s why you need manual transaction control.”

## How to Turn Off Autocommit

“In many systems, you can disable autocommit like this:

- In **MySQL**, run:

```
sql  
CopyEdit
```

```
SET autocommit = 0;
```

- 

- In **PostgreSQL**, start a transaction manually:

```
sql  
CopyEdit
```

```
BEGIN;
```

-

and it stays open until you COMMIT or ROLLBACK.

- In **programming languages** (like Python or Java), database drivers usually let you turn off autocommit in the connection settings.

Once autocommit is off:

- You control when to **commit**.
- You can **rollback** if there's a problem.
- You get the full power of transactions.”



### Key Message for Students

“Autocommit is good for quick and simple changes.

But for anything important or multi-step, **turn it off** so you can use proper transactions.”

## 5. Isolation Levels – Explained for Beginners



### What Are Isolation Levels?

“Imagine a busy kitchen where multiple chefs are preparing different meals. If they start using the same ingredients at the same time without rules, there'll be chaos — missing items, wrong orders, confusion.

The same happens in a **multi-user database**. If many people (or apps) are reading and writing at the same time, things can go wrong.

**Isolation levels** are the rules that control how transactions interact with each other. They decide how much one transaction can ‘see’ or be affected by others.”



### Why Isolation Matters

“Let's say:

- One transaction is reading the price of a product.
- Another transaction is updating that price at the same time.

Depending on the isolation level, the first transaction might:

- See the **old price**
- See the **new price**
- Or even see a price that's not finalized (a *dirty read*)!

That's why isolation levels exist: to **protect your data** when multiple things are happening at once."



## The Four Isolation Levels

Explain them **from weakest to strongest**:

### ◆ 1. Read Uncommitted

- **Most relaxed level.**
- Transactions can see changes made by others even if those changes **haven't been committed yet**.
- This allows **dirty reads** — risky because the other transaction might roll back.

"It's like reading someone's notes before they've finished writing. It might be wrong or get erased."

### ◆ 2. Read Committed

- Can only read data that's **already committed**.
- No dirty reads.
- But if another transaction updates the data in the meantime, you may see **different values if you read twice**.

"It's like checking a webpage twice and seeing different content each time."

### ◆ 3. Repeatable Read

- When you read data once, it **stays the same** for the whole transaction.
- Prevents **non-repeatable reads**.
- But **new rows** inserted by others may still be visible (phantom reads).

"It's like checking a document and freezing what you see — even if someone else edits it later."

### ◆ 4. Serializable

- **Most strict and safest.**
- Transactions are completely **isolated** — as if they were running one at a time.

- Prevents all kinds of concurrency problems: dirty reads, non-repeatable reads, phantom reads.

“It’s like you’re the only person in the library — no one else can change anything while you’re working.”

## How to Set Isolation Levels

“In most databases, you can set the isolation level at the beginning of a transaction:

```
sql
CopyEdit
SET TRANSACTION ISOLATION LEVEL READ COMMITTED;
BEGIN;
-- your SQL here
COMMIT;
```

You can also set it at the session level depending on your database.”

## Trade-offs: Performance vs. Safety

- **Lower isolation** (like Read Uncommitted) gives **better performance** — faster and fewer locks.
- **Higher isolation** (like Serializable) gives **more safety** — but can slow things down due to more locking and blocking.

“You have to **balance safety and speed**. For critical operations, use strict levels. For quick, non-critical reads, a lower level may be okay.”

## Key Message for Students

“Think of isolation levels as how *private* your transaction is. The more private (isolated) it is, the safer it is — but also slower. Choose the level based on how important accuracy is for your task.”

# 6. Common Transaction Problems – SQL Server Focus

## Why Transaction Problems Happen

“When we don’t carefully manage transactions — especially in systems with many users — problems can pop up. These problems don’t always show up right away, but they can lead to **wrong results, inconsistent reports**, and even **data corruption** over time.

Let’s go over the four major problems.”

## ◆ 1. Dirty Reads

### What is it?

A dirty read happens when a transaction reads data that was **changed by another transaction but not committed yet**. If that other transaction later rolls back, the data you read was never real — it was temporary and incorrect.

### Example:

One transaction updates a customer balance. Another transaction reads the new balance before it's committed. Later, the update is rolled back — but the second transaction already used that false balance.

### SQL Server Isolation Level to Prevent:

Use **Read Committed** or higher to block dirty reads.

(Default in SQL Server is `READ COMMITTED`, which **does block** dirty reads.)

## ◆ 2. Non-repeatable Reads

### What is it?

When a transaction reads the same row twice and **gets different values each time** because another transaction changed the row in between.

### Example:

You read a product's price. Another transaction changes it. You read it again and see a different price — in the same transaction.

### SQL Server Isolation Level to Prevent:

Use **Repeatable Read** or higher.

This level **locks** rows for reading, so no other transaction can change them until your transaction ends.

## ◆ 3. Phantom Reads

### What is it?

When a transaction reads a **set of rows** based on a condition, but later in the same transaction, it reads again and **sees new rows** that weren't there before — because another transaction inserted them.

### Example:

You run a query to get all orders over \$100. Another user adds a new \$150 order. When you run the same query again, the result has changed.

### SQL Server Isolation Level to Prevent:

Use **Serializable**.

It **locks the entire range of rows**, so others can't add new matching rows until your transaction is done.

## ◆ 4. Lost Updates

### What is it?

Two transactions **read the same data and update it**, but one of the updates is **overwritten** by the other without warning.

### Example:

Two people change the same product stock level at the same time. The last one to commit **overwrites** the changes made by the first.

### SQL Server Fix:

This needs **explicit row locking** or handling with isolation level like **Repeatable Read** or **optimistic concurrency** (checking if data changed before updating).

## 🧠 Key Summary for Students

Problem	Fix by Using Isolation Level
Dirty Reads	Read Committed or higher
Non-repeatable Reads	Repeatable Read or higher
Phantom Reads	Serializable only
Lost Updates	Careful locking or versioning

## 7. Transactions and Error Handling – SQL Server (T-SQL)

### ✅ Why Link Transactions with Error Handling?

“Even when we use transactions, things can go wrong:

- A statement might fail.
- A constraint might be violated.
- A deadlock might happen.

So we need to make sure that if something **goes wrong inside a transaction**, we don't accidentally **commit broken or partial data.**”

### 🔄 How to Handle Errors in Transactions – T-SQL Style

In **SQL Server**, we use **TRY . . . CATCH** blocks in T-SQL to handle errors inside transactions.



## ◆ Basic Structure:

```
sql
CopyEdit
BEGIN TRY
    BEGIN TRANSACTION;

    -- your SQL statements here

    COMMIT;
END TRY
BEGIN CATCH
    ROLLBACK;

    -- handle the error (optional logging, message, etc.)
    PRINT ERROR_MESSAGE();
END CATCH;
```

## 📌 Explain This Step-by-Step

1. **BEGIN TRANSACTION** – Start the transaction.
2. Inside the **TRY**, run your SQL statements.
3. If everything works, **COMMIT** saves the changes.
4. If anything fails, SQL Server jumps to the **CATCH** block.
5. In the **CATCH**, you **ROLLBACK** to undo all changes.
6. You can also log the error or show a message.

## 🧠 Why This Matters

“If we don’t catch errors and roll back, SQL Server might leave the transaction open, or worse, partially applied. That means **some changes happened, and some didn’t** — exactly what we don’t want.”

## 🧰 Best Practices for Students

- Always pair transactions with error handling.
- Use **TRY . . . CATCH** when doing inserts, updates, or deletes in groups.

- Don't forget the `ROLLBACK` in the `CATCH` block.
- Optionally, use `XACT_STATE ( )` to check if the transaction is still valid before committing.